

---

# Part 1 — Mark ECN

## Programming the Data Plane with DOCA Flow

*Programming SmartNICs: From Packet Processing to Programmable Transport*

ACM SIGCOMM 2026

Fernando Ramos  Muhammad Shahbaz  Mina Tahmasbi Arashloo  Mario Baldi 

Francisco Pereira  Bilal Saleem  Tyler Liu  Chenxingyu Zhao 

 INESC-ID, Instituto Superior Técnico, University of Lisbon  University of Michigan  University of Waterloo  
 NVIDIA  Purdue University  University of Washington

---

In this part you program the **data plane** of a BlueField-3: you tell the NIC what to do with packets *in its own hardware, at line rate*, before any CPU sees them. In this part of the tutorial, you will write a program that marks live RoCE traffic with an ECN congestion signal. Later, on the second part of the tutorial, you will build a congestion controller on the Bluefield that reacts to those same signals.

### Step 1 — The card, and getting traffic onto it

Your card has **two ports reachable to each other (p0 and p1)**, so whatever leaves p1 arrives at p0, and vice versa. That makes a single card behave like a small two-node network talking to itself. In this tutorial, we will be using two endpoints composed of lightweight virtual NICs called **SFs** (sub-functions), one per port, each in its own network namespace so the kernel cannot short-circuit them locally:

**INFO — what is a sub-function?** A **sub-function (SF)** is a lightweight virtual NIC carved out of a physical port. It shows up as its own network device. We create one SF on each side and use the two of them as the *endpoints* of a network flow — so a single card can play both “sender” and “receiver” across the cable.

The client sends RDMA WRITES over p1, and are later received by p0 and ultimately fed to the server.

What you will be reproducing is an everyday situation in a datacenter network: a sender is going as fast as it can, the network becomes congested, a switch on the path signals that congestion by setting the **ECN** bits in the IP header, and the sender slows down. Here, your card simulates the congested switch — no congestion actually exists, but you will be *manufacturing* the congestion signal by configuring the Bluefield to mark the ECN bits on the packets.

As such, we split this tutorial into two parts:

- **Part 1, this guide.** Program the NIC to set the ECN “congestion experienced” (CE) mark on a fraction of the client’s packets going through the NIC — you choose the fraction. Everything happens in hardware; your program only installs the rules, then sits idle printing counters.
- **Part 2.** Program the Bluefield’s DPA to *react*: the server’s NIC answers a CE-marked packet with a congestion notification, and your algorithm turns each one into a lower send rate for the client.

**The network devices of the physical ports.** The two physical ports are visible as soon as you log in:

```
$ ip -br link show | grep -E '^p0|^p1'
p0          UP      f0:fb:7f:e2:e2:76 <BROADCAST,MULTICAST,UP,LOWER_UP>
p1          UP      f0:fb:7f:e2:e2:77 <BROADCAST,MULTICAST,UP,LOWER_UP>
```

What the tutorial builds: a congestion signal, and something that reacts to it




Figure 1: You write the marking in Part 1 and the reaction in Part 2. The client, the server, and the CNP the receiver's NIC sends back are already there.

These MAC addresses are specific to whichever card you are using, so yours might read differently.

**The RoCE endpoints.** In this tutorial, we will work with **RoCE** (RDMA over Converged Ethernet) traffic, the high-speed, kernel-bypass transport used in AI and storage networks. RoCE doesn't use normal sockets. Instead, programs reach it through **RDMA devices** named `mlx5_{N}`. Each SF has one, and it is only visible from inside that SF's namespace:

```
$ sudo ip netns exec ns0 rdma link show
link mlx5_2/1 state ACTIVE physical_state LINK_UP netdev enp3s0f0s0 # the server's SF
$ sudo ip netns exec ns1 rdma link show
link mlx5_3/1 state ACTIVE physical_state LINK_UP netdev enp3s0f1s0 # the client's SF
```

**Run `rdma link show` outside a namespace and none of that appears.** You get `mlx5_0` and `mlx5_1` — the two physical ports — each listing dozens of ports, most of them DOWN. That is the switch's own view, not the endpoints'. `mlx5_2` and `mlx5_3` are the names you want, and `ip netns exec` is what makes them visible.

**Step 1.1: wire up the two endpoints.** The two ports are wired into a loopback and split into two isolated network sandboxes (Linux network namespaces) called `ns0` and `ns1`, one SF in each, each with its own IP address (`mlx5_2` → `ns0` → `10.0.0.1`, `mlx5_3` → `ns1` → `10.0.0.2`). This is already set up for you on the tutorial card, so *there's nothing to run here*.

| Endpoint     | RDMA device         | Namespace        | IP                    | Role                          |
|--------------|---------------------|------------------|-----------------------|-------------------------------|
| SF on port 0 | <code>mlx5_2</code> | <code>ns0</code> | <code>10.0.0.1</code> | <b>RoCE server</b> (receives) |
| SF on port 1 | <code>mlx5_3</code> | <code>ns1</code> | <code>10.0.0.2</code> | <b>RoCE client</b> (sends)    |

**INFO — what is a network namespace?** It is a private, isolated network stack inside one Linux machine — its own interfaces, IPs, and routes. Putting each SF in its own namespace (`ns0`, `ns1`) is what makes them behave like two separate hosts even though they live on the same card. You run a command “inside” a namespace with `ip netns exec <name> <command>`.

**Step 1.2: put real traffic on the loopback.** We generate traffic with `ib_write_bw` — a standard RoCE benchmarking tool (from the `perftest` package) that measures how fast one endpoint can write data to another. It needs a server (receiver) and a client (sender).

**Try it yourself! Send RoCE across the cable and measure it!**

Open **two terminals**, both on the Arm cores.

**Terminal 1 — the receiver (server).** Start this one first; it waits for a client:

```
sudo ip netns exec ns0 ib_write_bw -d mlx5_2 -R -x 1 -F --report_gbits
```

**INFO - what the flags mean:** `ip netns exec ns0` runs it inside the `ns0` sandbox; `-d mlx5_2` uses that namespace's RDMA device; `-R` sets the connection up via the RDMA connection manager (keep this on — Part II needs it); `-x 1` picks the RoCEv2 address; `-F` ignores a CPU-frequency warning; `--report_gbits` prints the result in gigabits/second. It prints its settings and then says it is **waiting for a client**.

**Terminal 2 — the sender (client).** Point it at the server's IP, 10.0.0.1:

```
sudo ip netns exec ns1 ib_write_bw -d mlx5_3 -R -x 1 -F 10.0.0.1 --report_gbits
```

**You should see** a results table appear on both terminals, with the throughput closely matching the card's line rate:

| #bytes | #iterations | BW peak[Gb/sec] | BW average[Gb/sec] | MsgRate[Mpps] |
|--------|-------------|-----------------|--------------------|---------------|
| 65536  | 5000        | 92.58           | 92.56              | 0.176543      |

To avoid retyping the flags, the repo wraps these as scripts, in `scripts/` at the top of the repository: `run_server.sh` and `run_client.sh` (one per terminal), or **`benchmark.sh`**, which starts both ends together in a single command and streams the sender's throughput (Ctrl-C stops both). We recommend using `benchmark.sh` from Step 3 onwards, where it is reached as `../scripts/benchmark.sh`.

## Step 2 — Understanding DOCA Flow

Interacting with a BlueField-3 typically entails interfacing with — and often independently programming — four different architectural components:

| Where                     | What it is                               | Used for                     |
|---------------------------|------------------------------------------|------------------------------|
| <b>Host x86</b>           | The server the card is plugged into      | Not used here                |
| <b>Arm cores</b>          | General purpose cores on the card        | Where you'll be working      |
| <b>NIC ASIC / eSwitch</b> | The match-action pipeline inside the NIC | Packet modification (Part 1) |
| <b>DPA</b>                | RISC-V engine on the data path           | Congestion control (Part 2)  |

DOCA Flow is the NVIDIA library that allows you to program the Bluefield's eSwitch, allowing you to make packet modifications on the fly as they traverse the NIC's ASIC. Your DOCA Flow program builds a small graph of **pipes**, and each pipe answers four questions:

| Part           | The question it answers                     | Example                                |
|----------------|---------------------------------------------|----------------------------------------|
| <b>match</b>   | <i>which</i> packets does this pipe act on? | all IPv4 packets                       |
| <b>monitor</b> | <i>how many</i> packets hit it?             | a hardware counter you can read        |
| <b>actions</b> | <i>what</i> do we change in the packet?     | modify the destination IP address      |
| <b>forward</b> | <i>where</i> does the packet go next?       | another pipe, a port, the CPU, or drop |

You describe these rules once, at startup. The NIC then applies them to every packet in hardware, while your program sits idle printing counters. That is the whole idea: **match, count, modify, forward**.

Two details about pipes are worth getting straight before you write any:

- **Pipe = which fields. Entry = what values.** A pipe is a *template*: writing `0xFF` into a field of its match means “this field participates”, and `0x00` in the mask means “...but do not actually compare




Figure 2: Pipes chain into a graph. Within a pipe, each entry is a Match, a monitor (MON), a modify/action stage (MDF) and a forward (FWD). Source: [NVIDIA DOCA Flow programming guide](#), “Architecture”.

it”. The **entry** you add afterwards supplies the value really compared. Creating a pipe installs no rule in the hardware; adding an entry does. The same split applies to actions — the pipe declares “entries may rewrite this field”, the entry says “...to this value”.

- **One pipe is the root.** Every packet entering the eSwitch starts its lookup at the pipe marked `is_root`. No packet reaches other pipes unless the root sends it there. The program you are given ships with a root pipe that forwards packets to the server; the exercise is to put your own pipeline in between.

### Step 3 — Build the pipeline

The repository (`/home/s26t/sigcomm26-tutorial-bluefield-participants`) has one directory per DOCA release: `doca-2` and `doca-3`. Pick the relevant one for your chosen Bluefield.

**Which of the two is yours?** Run this command to find out which version of DOCA is installed on your card:

```
$ cat /opt/mellanox/doca/applications/VERSION
2.9.1008
```

Only the first number matters: **2.x → use doca-2**, **3.x → use doca-3**. The exercises are identical in the two versions, down to the same three TODOs, differing only in a few DOCA calls renamed between the generations, which are already written for you in each tree.

**Move into your version’s directory now, and stay there for the rest of Part 1:**

```
$ cd ~/sigcomm26-tutorial-bluefield-participants/doca-2 # or doca-3
```

This is the only place the version is ever named. Every command from here on is written relative to that directory, so you can paste them as they are whichever card you are on.

In this part of the tutorial, you will be editing the `doca-flow/doca_flow_ecn.c` file. Everything in `doca-flow/doca_flow_ecn.c` is done except for three function bodies, marked TODO 1 to TODO 3. The

program already forwards traffic as shipped.

Laid out logically, the `doca_flow_ecn.c` file is logically separated into four different phases. We use python-like pseudo-code to show you these logical components:

```
def doca_flow_ecn_main():
    # Phase 1: setting up logging and argument parsing
    setup_logging()
    parse_args()                # --percent

    # Phase 2: bring-up the device and library. None of this is pipeline work
    dev = open_and_probe_dev(0)  # PF0, probed into DPDK with its SF representor
    configure_and_start_dpdk_port(dev) # packet buffer pool, RX/TX queues
    initialize_doca_flow()        # switch mode, hardware steering, counters
    port = port_start(dev)        # the PF0 proxy port -- must be started first
    sf = rep_port_start(...)      # the server SF's representor

    # Phase 3: actually program the DOCA Flow pipeline
    build_pipeline(port, cfg, out) # <-- **all** of your work will be here

    # Phase 4: runtime logic (just reporting) and teardown
    run_report_loop()            # print the counters, once a second
    teardown()
```

Those are the real function names, so you can grep for any of them. The pipeline itself is six functions at the bottom of the file:

| Function                                 |               | What it does                                      |
|------------------------------------------|---------------|---------------------------------------------------|
| <code>build_pipeline()</code>            | given         | Calls the pipe creation functions                 |
| <code>create_passthrough_pipe()</code>   | given         | Forward everything to the server, unchanged       |
| <code>create_root_pipe_nop()</code>      | given         | The root pipe as shipped: wire to server and back |
| <code>create_root_pipe()</code>          | <b>TODO 1</b> | Yours — the same, but into your chain             |
| <code>create_forward_to_sf_pipe()</code> | <b>TODO 2</b> | Forward, count, and set the CE mark when asked    |
| <code>create_sampling_pipe()</code>      | <b>TODO 3</b> | Send 1 packet in N down the marking path          |

`build_pipeline()` is the one to read closely: it is the whole exercise in a single function, and it ships wired as a plain forwarder. In this function, note what `wire_target` does: it starts out as `PASSTHROUGH` — a working forwarder on its own, and all Step 4.1 needs — and is then reassigned to whichever pipe wire traffic should really enter as you add the pipes in front of it. `PASSTHROUGH` stays in the pipeline either way: it is where `MARK` and `PASS` send anything they do not match.

## Step 4 — The actual tutorial exercise

We use `meson` and `ninja` to setup and build our applications. You only need to setup once, but you do need to recompile with `ninja` every time you make changes to your program. All three commands run from inside the version directory you moved into in Step 3:

```
# Setting up the build directory (only once)
$ meson setup build

# Compile the application
$ ninja -C build

# Run the doca flow application
$ sudo ./build/doca-flow/doca_flow_ecn -- --percent 100
```

**The -- matters.** Everything before it goes to the DPDK library; everything after it is for this program. --percent N is the only flag it takes: CE-mark this share of packets, [0, 100], default 100.

**Keep ../scripts/benchmark.sh running in a second shell** throughout. It starts the server and the client together and streams the sender's throughput, so you can see the effect of every change.

**Figure 3 shows the DOCA Flow pipeline you will build.** Each blue pipe is one function, written in the section it is labelled with. Not drawn: PORT\_DEMUX's second entry, sending everything coming back from the server straight out to the wire.




Figure 3: The Step 4 pipeline. Solid arrows are where a packet goes when it matches, dashed ones where it goes when it does not.

### Step 4.1 — Write the root pipe (PORT\_DEMUX)

**First, run it exactly as it comes.** Traffic should sit at line rate — 92 or 184 Gb/s depending on the card — and the counter line should stay at zero:

```

s26t@bluefield-lisbon-1:~/sigcomm26-tutorial-bluefield-participants/doca-2$ sudo
↪ ./build/doca-flow/doca_flow_ecn -- --percent 100
EAL: Detected CPU lcores: 16
EAL: Detected NUMA nodes: 1
EAL: Detected shared linkage of DPDK
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'VA'
TELEMETRY: No legacy callbacks, legacy socket not created
EAL: Probe PCI driver: mlx5_pci (15b3:a2dc) device: 0000:03:00.0 (socket -1)
[01:49:05:767697][2777301][DOCA][INF][doca_flow_ecn.c:213][configure_and_start_dpdk_port] mbuf
↪ data room 9344 bytes (jumbo-capable, max frame 9216)
[01:49:06:723850][2777301][DOCA][WRN][engine_model.c:92][adapt_queue_depth] adapting queue
↪ depth to 128.
[01:49:07:756283][2777301][DOCA][INF][doca_flow_ecn.c:693][create_root_pipe_nop] No-op
↪ forwarder ready: wire <-> receiver SF, nothing marked
[01:49:07:756313][2777301][DOCA][INF][doca_flow_ecn.c:383][log_startup] Marking ALL IPv4 --
↪ Ctrl-C to stop
[01:49:08:004432][2777301][DOCA][INF][doca_flow_ecn.c:401][run_report_loop] CE marked: 0,
↪ passthrough: 0 (0% marked)
[01:49:09:004436][2777301][DOCA][INF][doca_flow_ecn.c:401][run_report_loop] CE marked: 0,
↪ passthrough: 0 (0% marked)
  
```

Currently, the pipeline is pre-configured with a root pipe that simply forwards packets from the wire straight to the server, and performing no additional operation. This pipe is created in the function `create_root_pipe_nop()`.

**Now stop it with Ctrl-C, leaving the traffic (benchmark.sh) running.** Throughput carries on unchanged, aside from a small dip in throughput as the card falls back to its default OVS forwarding. Start

the program again and it takes over (notice again the small temporary dip in throughput).

**Why this matters.** The instant a DOCA Flow program starts, it **owns the NIC's switch**, and from then on the NIC forwards only what your pipes say to forward. `create_root_pipe_nop()` is what keeps traffic moving. Take it away with nothing in its place and everything stops: an empty pipeline does not mean “pass traffic through”, it means “drop everything”.

**Updating the pipeline building mechanism.** In `build_pipeline()`, comment out the call to the no-op root pipe (`create_root_pipe_nop()`), and uncomment the two lines tagged `// [1]`.

**Now start writing your own root pipe.** Let's now fill in `create_root_pipe()` — its two gaps are `TODO 1a` and `TODO 1b`. Every pipe is built the same way, in two layers:

- **The pipe is a template.** When you create a pipe you describe its *shape* — which header fields it matches on, which fields its actions may rewrite, and where a match forwards — but **not** the concrete values. You mark a field with `0xFF` to mean “this field participates,” without yet saying *what* to look for. The call that compiles that shape into the hardware is `doca_flow_pipe_create()`.
- **Entries fill in the values.** A freshly created pipe does nothing until you add at least one **entry**, and each entry supplies the concrete values for the fields the pipe declared — “match IPv4 packets *with this DSCP/ECN byte*,” “rewrite the ECN bits *to this value*.” One pipe can carry many entries (a root pipe, for instance, has one per direction). You add an entry with `doca_flow_pipe_add_entry()`, then install it in the NIC with `doca_flow_entries_process()`.

So the one idiom behind every pipe: **the pipe says *which fields*, the entry says *what values*** — and the same split applies to actions (the pipe declares “entries may rewrite this field,” the entry supplies the value to write).

You are not writing from a blank page. Take heavy inspirations from `create_root_pipe_nop()` directly above `create_root_pipe()`, because you are writing almost the same pipe:

- **(TODO 1a) Match** on the ingress port, `match.parser_meta`, with a full mask so it is compared exactly. This is the field that says which side a packet came from.
- **(TODO 1a) Forwards:** `DOCA_FLOW_FWD_CHANGEABLE` for a hit — meaning “each entry brings its own destination” — and `DOCA_FLOW_FWD_DROP` for a miss.
- **(TODO 1b) Two entries.** From the wire (`PF_PORT_ID`) to `wire_target`, and from the server's SF (`SF_REP_PORT_ID`) out to `PF_PORT_ID`. Install the first with `WAIT_FOR_BATCH` and the second with `NO_WAIT`, so both reach the hardware together.

**The only real difference from the no-op** is that first entry: the no-op forwards wire traffic to a **port** (`DOCA_FLOW_FWD_PORT`), yours forwards it to a **pipe** (`DOCA_FLOW_FWD_PIPE`, `next_pipe = wire_target`). That is the whole distinction between a forwarder and a pipeline, and it is what lets you insert anything at all into the path.

Keeping the two directions apart is not cosmetic. The return path carries the RoCE acknowledgements and the congestion notifications the Part 2 controller reacts to; marking those would corrupt the very feedback you are trying to create.

Rebuild and run. Traffic should be back at line rate, now through your pipe rather than the shipped one, with the counters still at zero — `wire_target` is `PASSTHROUGH`, which counts nothing.

## Step 4.2 — Mark every packet

Uncomment the rest of the block in `build_pipeline()`, then fill in `create_forward_to_sf_pipe()` — `TODO 2a` and `TODO 2b`.

**This one function builds both of the forwarding pipes.** `build_pipeline()` calls it **twice**: once to build a pipe that does *not* mark the packets' ECN bits (“PASS” in Figure 3), and another to build a pipe that does mark them (“MARK” in Figure 3). Everything the two have in common — the match, the counter,



the forwards — you write unconditionally. The two steps that are specific to the marking pipe are the action template in TODO 2a and the value the entry writes in TODO 2b.

Take inspiration from `create_passthrough_pipe()`, right above `create_forward_to_sf_pipe()`: this is essentially the same pipe but with a counter and an action added.

- **Match** any IPv4 packet *whatever ECN bits it arrived with*, using the wildcard idiom from Step 2: `outer.ip4.dscp_ecn` as `0xFF` in the pipe's match, `0x00` in the mask. Reset that byte to `0x00` before adding the entry — the same struct is reused as the entry's values.
- **A counter**, via `monitor.counter_type = DOCA_FLOW_RESOURCE_TYPE_NON_SHARED`. Without it the CE marked: line stays at zero and you cannot tell whether anything works.
- **The marking action**, when `mark` is true: declare `outer.ip4.dscp_ecn` as `0xFF` in a pipe-level action template, and have the entry write the value. RFC 3168 gives the two-bit ECN field as Not-ECT 00, ECT(1) 01, ECT(0) 10, CE 11, so the byte you want is `0x03`.
- **Forwards**: a hit goes to the server's SF; a miss goes to `miss_pipe`, which is `PASSTHROUGH` and forwards everything. Nothing is dropped here — the match only decides what gets counted and marked, so non-IPv4 (ARP and the like) travels on untouched. A miss may only be a pipe or a drop, never a port, which is why the fallback is a pipe.
- Hand the installed entry back through `out_entry` — that is what the counter report queries.

Rebuild and run with `--percent 100`. Traffic should be at line rate, and now the counter climbs:

```
CE marked: 69144890, passthrough: 0 (100% marked)
CE marked: 80447146, passthrough: 0 (100% marked)
```

**You just programmed the Bluefield to rewrite packet headers in hardware**, at line rate, with your program running on the Arm cores doing nothing but printing the counter once a second.

### Step 4.3 — Mark only some packets

We will now create Figure 3's `RANDOM_SAMPLE` pipe by implementing `create_sampling_pipe()` (TODO 3a and TODO 3b), which splits traffic across pipes probabilistically in hardware.

The NIC stamps every packet with a random 16-bit value in `parser_meta.random`. Match it against 0 under mask, which has already been computed for you as a power of two minus one, and exactly 1 packet in `(mask + 1)` hits. Hits forward to the marking pipe, misses to the non-marking one; “miss” here means *not selected*, not an error. The entry itself adds nothing to the template — no actions, no counter, no forward of its own.

```
$ sudo ./build/doca-flow/doca_flow_ecn -- --percent 12.5
```

The startup banner prints the fraction actually achieved, rounded down to a power of two. Check it against the counter line, which now has traffic on both sides:

```
CE marked: 39273726, passthrough: 274903125 (12.5% marked)
CE marked: 40684757, passthrough: 284794849 (12.5% marked)
```

### Checking your work

There is no answer key in your checkout, and you do not need one: the program tells you where you stand at every step. Traffic back at line rate ends Step 4.1, CE marked: climbing ends Step 4.2, and the two counters splitting in the ratio you asked for ends Step 4.3. When one of those does not happen, the debugging tips below name the usual cause.

Ask an organiser if you are stuck. That is what we are here for.



## Debugging tips

- **Read the *last* error line, not the first.** A failed pipe prints a wall of internal DOCA errors. The final `[CRT]...[doca_check] <name>: <reason>` line names the pipe, and that name is the function to open.
- **Nothing forwards at all** — expected in the middle of Step 4.1, once the no-op call is commented out and `create_root_pipe()` is still empty. If it persists, check that `doca_flow_pipe_create()` ran and that `doca_flow_entries_process()` accounted for both entries.
- **Traffic stops the moment you uncomment the rest of the block** — a forward points at a pipe that was never built. `create_forward_to_sf_pipe()` and `create_sampling_pipe()` return NULL until you write them, and a root pipe aimed at NULL forwards nowhere.
- **CE marked: stays at 0** — your marking pipe has no counter (`monitor.counter_type`), or the root pipe is still aimed at PASSTHROUGH rather than at the marking pipe.
- **The client connects and then stalls** — the root pipe's second entry, server SF back out to the wire, is missing or points the wrong way. RoCE needs both directions.
- **A pipe fails to install** — check the status after `doca_flow_entries_process()`, per the pipe-and-entry idiom in Step 4.1. Success from the add-entry call alone proves nothing.
- **EAL initialization failed, then argp: doca\_argp\_start(...)** — a copy of the program is already running, and only one DOCA Flow program can own the switch at a time. This is the one case where the last line misleads: it blames argp, but nothing is wrong with your arguments. Stop the other instance; if it was killed hard, `sudo pkill -f doca_flow`.
- **defined but not used warnings** — that is the list of functions you have not wired up yet. Expect five before you start; they clear as you uncomment in Step 4.1 and Step 4.2.

## What you built

You can now treat the card as a two-port loopback carrying real RoCE traffic; follow how match/count/modify/forward pipes chain from a root pipe; and write a pipeline that takes over the NIC's switch, matches IPv4, counts it, rewrites its ECN bits to CE, and forwards it — all in hardware, at line rate, with your program idle. That CE mark is exactly the signal the congestion controller in Part 2 reacts to.

## Appendix A — The BlueField, in more detail

*Reference material. You do not need any of this to finish the exercise.*

Interacting with a BlueField-3 typically entails interfacing with — and often independently programming — four different architectural components:

| Where                     | What it is                                                           | Used for                          |
|---------------------------|----------------------------------------------------------------------|-----------------------------------|
| <b>Host x86</b>           | The server the card is plugged into                                  | Not used at all in this tutorial  |
| <b>Arm cores</b>          | Cortex-A78 cores running their own Ubuntu on the card                | Where you log in, build and run   |
| <b>NIC ASIC / eSwitch</b> | The match-action switch inside the NIC                               | What you program in Part 1        |
| <b>DPA</b>                | Datapath Accelerator: a many-threaded RISC-V engine on the data path | The congestion controller, Part 2 |

The cards are configured in **DPU mode** (also called embedded-CPU-function ownership, ECPF): the Arm subsystem owns the NIC's resources, and the host x86 sees only a plain network device.

### Functions: PFs, VFs and SFs

A **function** is what the NIC exposes so that software can send and receive packets through it. It is not a piece of software, but a slice of the NIC hardware, with its own queues, its own MAC address, and its own identity on the network. There are three kinds of functions:

- **PF** (*physical function*): the real PCIe device the NIC presents. A BlueField-3 has one PF per physical port. These are the functions that own the hardware.
- **VF** (*virtual function*): an [SR-IOV](#) slice of a PF, created so that a virtual machine can be handed something that looks like its own NIC. Not used in this tutorial.
- **SF** (*scalable function*, also called a *sub-function*): the same idea as a VF, but not tied to PCIe. An SF is created on the Arm with `m1nx-sf`, is much cheaper than a VF, and you can have far more of them. **This tutorial uses SFs** as its traffic endpoints.

Each function shows up on the Arm's Linux as **two** different devices:

| Interface          | What it is                                                                                | Example             |
|--------------------|-------------------------------------------------------------------------------------------|---------------------|
| <b>netdev</b>      | An ordinary Linux interface — what <code>ip link</code> lists and <code>ping</code> uses. | <code>p0</code>     |
| <b>RDMA device</b> | The same function through the RDMA stack, bypassing the kernel.                           | <code>m1x5_0</code> |

In this tutorial we use RDMA over Converged Ethernet (**RoCE**), which is RDMA carried inside ordinary UDP/IP packets, so we use the RDMA devices (`m1x5_N`). Because RoCE is just UDP/IP on the wire, the switch inside the NIC can match and rewrite it like any other packet.

| Function  | RDMA device         | Netdev                  |
|-----------|---------------------|-------------------------|
| PF0       | <code>m1x5_0</code> | <code>p0</code>         |
| PF1       | <code>m1x5_1</code> | <code>p1</code>         |
| SF on PF0 | <code>m1x5_2</code> | <code>enp3s0f0s0</code> |
| SF on PF1 | <code>m1x5_3</code> | <code>enp3s0f1s0</code> |

### The eSwitch

The eSwitch (embedded switch) is a hardware block inside the NIC that forwards packets between every function and every physical port on the card. Every packet that arrives from the wire and every packet

a function transmits passes through it, and it decides where that packet goes. It is the single most important component in this part of the tutorial, and we use DOCA Flow to program it.

Its rule table is called the **FDB** (forwarding database). Each PF (port) owns its own logical eSwitch domain, with its own rules, so programming PF0’s eSwitch leaves PF1’s forwarding untouched.

Every endpoint the eSwitch can send a packet to is a **vport** (virtual port): the physical ports are vports, and so is every PF, VF and SF. The eSwitch does not reach an SF directly, though. It reaches the SF’s **representor** — a switch-side netdev that stands in for it. The two are the same function seen from opposite sides:

- from *inside* the namespace or VM using it, the SF appears as its netdev and its `m1x5_N` RDMA device — this is where an application sends and receives;
- from the *switch*’s side, the same SF appears as its representor, which is what a forwarding rule names.

So “deliver this packet to the receiver” is written, in a rule, as “output to that representor’s vport” — and the packet then pops out of the corresponding `m1x5_N` device inside the namespace. In NVIDIA’s figure below, `rep0` and `rep1` are the representors of `VF0` and `VF1`; substitute SFs for VFs and it is our layout exactly, with the DOCA Flow application running on the Arm.




Figure 4: Applications use the functions (`VF0`, `VF1`); forwarding rules name their representors (`rep0`, `rep1`) and the physical ports. Source: [NVIDIA DOCA Flow programming guide](#), “Domains in Switch Mode”.

By default, each PF’s vports are wired together by an OVS bridge with hardware offload — the “normal” forwarding path that makes the card behave like an ordinary NIC when nothing else is running. A DOCA Flow program **replaces** that for the PF it attaches to, from the moment it starts until it exits.

## The whole round trip

The client issues an RDMA WRITE through `m1x5_3`; PF1’s eSwitch — untouched by you — sends it out `p1`; it crosses the cable to `p0`; **PF0’s eSwitch, your pipeline, marks and forwards it**; the server receives it via `m1x5_2`; the server’s NIC answers a CE-marked packet with a **CNP** (Congestion Notification Packet) in hardware; the CNP travels back the same way; and on the client’s side it raises an event on the DPA, where the Part 2 algorithm sets a new send rate.

## Appendix B — DOCA Flow concepts in full

*Reference material. Step 2 has the working subset.*

**Port.** A DOCA Flow handle on one vport. In *switch mode* the PF uplink is also the **proxy port** (switch-manager port): it must be started first, and every pipe belongs to it — even pipes that forward somewhere else entirely.

DOCA Flow + DOCA PCC (USER\_PROGRAMMABLE\_CC=1, PCC loaded)




Figure 5: The end-to-end data path with both parts running.

**Pipe.** A flow table, not a pipeline stage. NVIDIA’s wording: a pipe “is a template that defines packet processing without adding any specific hardware rule”.

| Pipe type    | What it does                                                                      |
|--------------|-----------------------------------------------------------------------------------|
| BASIC        | Ordinary match-action table                                                       |
| CONTROL      | Entries carry priorities (0-7), resolving conflicts between overlapping matches   |
| LPM          | Longest-prefix match, for routing-table-shaped lookups                            |
| ACL          | Five-tuple matching with masks                                                    |
| ORDERED_LIST | A fixed sequence of actions per entry, when their order matters                   |
| HASH         | Selects an entry by an index computed from a hash, rather than by matching fields |

**Entry.** A concrete rule inside a pipe.

**Match.** Each field is *ignored* (left zero), *constant* (set in the pipe, the same for all entries), or *changeable* (all-ones placeholder in the pipe, real value supplied per entry).

**Actions.** Mostly field rewrites — MAC, IP, DSCP/ECN, L4 ports, metadata — but also encapsulation and decapsulation. Matches and actions are not alternatives to one another: every entry has a match, and *may additionally* have actions, a monitor and a forward.

**Monitor.** Counters, meters, and (on 2.x) mirroring.

**Forward.** Where the packet goes when the lookup finishes. Each pipe has a *hit* forward and optionally a *miss* forward:

| Forward                  | Meaning                                                             |
|--------------------------|---------------------------------------------------------------------|
| DOCA_FLOW_FWD_PORT       | Output to a vport: a physical port, or a function’s representor     |
| DOCA_FLOW_FWD_PIPE       | Jump to another pipe — this is how pipes chain                      |
| DOCA_FLOW_FWD_RSS        | Deliver to a receive queue of your program, <b>on the Arm cores</b> |
| DOCA_FLOW_FWD_DROP       | Discard                                                             |
| DOCA_FLOW_FWD_CHANGEABLE | Each entry brings its own forward                                   |
| DOCA_FLOW_FWD_HASH_PIPE  | Jump to a hash pipe, naming the algorithm to run (3.x)              |

**Shared resources.** Objects living on the port rather than inside one pipe, so several pipes can point at a single instance: meters, counters, RSS contexts, encap/decap contexts — and, on 2.x, **mirrors**, which duplicate a packet towards a second destination. Mirrors were removed in DOCA 3.2, and copying traffic

into a capture file is where the two versions diverge hardest over it: a shared mirror on 2.x, a flooding hash pipe on 3.x. Neither is part of this exercise.

**Parser metadata.** `parser_meta` holds values the hardware parser attaches to each packet, rather than header fields: `outer_l3_type` (what the parser saw), `random` (a fresh 16-bit value per packet, which is what makes hardware sampling possible), and the ingress port.

Note that `outer.ip4.dscp_ecn` is the **whole** TOS byte, so writing `0x03` sets ECN to CE *and* clears DSCP. Our traffic carries DSCP 0, so it makes no difference here.

## Appendix C — The API calls used in this exercise

*The headers on the card are the authority: `/opt/mellanox/doca/include/doca_flow.h`, with packet field types in `doca_flow_net.h`. Every call is documented there. This is only a map of which ones matter here.*

`grep -n 'dscp_ecn\|parser_meta' /opt/mellanox/doca/include/doca_flow*.h` answers most structural questions faster than the web documentation. If you are on VS Code Remote-SSH, `.vscode/c_cpp_properties.json` already points IntelliSense at these paths, so “go to definition” works.

| Call                                              | What it is for                                                           |
|---------------------------------------------------|--------------------------------------------------------------------------|
| <code>doca_flow_pipe_cfg_create / _destroy</code> | Start and finish describing a pipe                                       |
| <code>doca_flow_pipe_cfg_set_name</code>          | Name it — this is what error messages report                             |
| <code>doca_flow_pipe_cfg_set_type</code>          | BASIC, HASH, ...                                                         |
| <code>doca_flow_pipe_cfg_set_domain</code>        | Which steering domain; always <code>..._DOMAIN_DEFAULT</code> here       |
| <code>doca_flow_pipe_cfg_set_is_root</code>       | Mark the one root pipe                                                   |
| <code>doca_flow_pipe_cfg_set_nr_entries</code>    | How many entries you will add                                            |
| <code>doca_flow_pipe_cfg_set_match</code>         | The match template and its mask                                          |
| <code>doca_flow_pipe_cfg_set_actions</code>       | The action template(s)                                                   |
| <code>doca_flow_pipe_cfg_set_monitor</code>       | The counter                                                              |
| <code>doca_flow_pipe_create</code>                | Create the pipe, with its hit and miss forwards                          |
| <code>doca_flow_pipe_add_entry</code>             | Add an entry — [2.x] and 3.1                                             |
| <code>doca_flow_pipe_basic_add_entry</code>       | Add an entry — [3.x]; takes the action index as an argument              |
| <code>doca_flow_entries_process</code>            | Drive queued entries to completion, then check the status                |
| <code>doca_flow_resource_query_entry</code>       | Read an entry’s counter (already written, in <code>query_pkts()</code> ) |

## Version differences

|              | [2.x]                                                                     | [3.x]                                                                           |
|--------------|---------------------------------------------------------------------------|---------------------------------------------------------------------------------|
| entry        | <code>doca_flow_pipe_add_entry</code> ; action index                      | <code>doca_flow_pipe_basic_add_entry</code> ;                                   |
| install      | inside <code>doca_flow_actions</code>                                     | action index as an argument                                                     |
| entry flags  | <code>DOCA_FLOW_NO_WAIT</code> ,<br><code>DOCA_FLOW_WAIT_FOR_BATCH</code> | <code>DOCA_FLOW_ENTRY_FLAGS_NO_WAIT</code> ,<br><code>..._WAIT_FOR_BATCH</code> |
| ingress port | <code>parser_meta.port_meta</code> (uint32)                               | <code>parser_meta.port_id</code> (uint16)                                       |
| match        |                                                                           |                                                                                 |

The version you are on is already handled for you: `doca_flow_ecn.c` in `doca-2/` uses the 2.x spelling throughout and the one in `doca-3/` uses the 3.x spelling, so follow whichever file you have open rather than translating between them.

`doca_flow_compat.h`, force-included by the build, already smooths the 3.1-versus-3.4 seam inside the `doca-3` tree. It does **not** bridge 2.x to 3.x.

## The program's own flags

| Flag      | Meaning                                                                              |
|-----------|--------------------------------------------------------------------------------------|
| --percent | CE-mark this share of packets, $[0, 100]$ , rounded down to a power-of-two fraction. |
| N         | Default 100.                                                                         |

## Further reading

- [DOCA Flow programming guide](#) — swap the version in the URL to match your card, e.g. .../archive/2-9-0/doca-flow.
- [BlueField modes of operation](#)
- [BlueField scalable function user guide](#)
- NVIDIA ships around 20 single-concept sample programs on the card. They live under /opt/mellanox/doca/samples/doca\_flow/ — one .c file plus a README each, every one isolating a single idea.