
Part 2 — Programmable Congestion Control with DOCA PCC

Programming SmartNICs: From Packet Processing to Programmable Transport

ACM SIGCOMM 2026

Fernando Ramos  Muhammad Shahbaz  Mina Tahmasbi Arashloo  Mario Baldi 

Francisco Pereira  Bilal Saleem   Tyler Liu  Chenxingyu Zhao 

 INESC-ID, Instituto Superior Técnico, University of Lisbon  University of Michigan  University of Waterloo
 NVIDIA  Purdue University  University of Washington

In **Part 1** ([doca-flow.pdf](#)) you programmed the NIC to *mark* congestion — you set the CE bit on packets in the data plane. In this part you write the other side of that story: the **congestion-control algorithm that reacts to those marks**, deciding how fast a flow is allowed to send — and you run it on a set of tiny processors *inside* the NIC called the **DPA**.

By the end you will have **written the two reactions at the heart of a reaction-point controller**, watched a flow’s send rate **collapse** the moment congestion appears and **recover** when it clears, and tuned how aggressively it reacts.

We build up in three parts:

- **Part A** — *where* the algorithm runs: the two halves (a host program that loads it, and the DPA that runs it), and how a “rate” is expressed.
- **Part B** — *how* the controller thinks: the event loop and the two reactions that make up a DCQCN-style loop — both of which you’ll write.
- **Part C** — fill in the two reactions one at a time: **TODO 1** (the cut) and watch the rate collapse, then **TODO 2** (the recovery) and watch the sawtooth; then tune how hard it reacts.

Prerequisites. - You have done **Part 1** ([doca-flow.pdf](#)) — PCC reacts to the ECN marks you produced there. - You are on the **Arm cores** of a BlueField-3 with this repo checked out and sudo access. - One firmware knob, `USER_PROGRAMMABLE_CC=1`, must be live for any PCC program to start — on the tutorial cards this is already set for you. (Details in the [FAQs](#).)

Part A — Where the algorithm runs

Two halves: a host loader and a DPA program

A DOCA PCC program comes in two pieces that run in two different places:

- **The host program** ([doca_pcc_ecn_rp](#)) is just a **loader and supervisor**. It uploads your compiled algorithm to the NIC, opens the PCC context, and then sits there keeping it alive and printing logs. It does **not** run the algorithm itself.
- **The algorithm** runs on the **DPA** — the *Data-Path Accelerator*, a cluster of small, highly parallel processors inside the BlueField-3. That is where the C code you’ll look at actually executes.

INFO — what is the DPA, and why not the Arm cores? The Arm cores are general-purpose Linux CPUs; they are too far from the wire to make a per-flow rate decision fast enough. The DPA sits right on the data path and is built for exactly this kind of tiny, frequent,




Figure 1: The two halves of a DOCA PCC program. The host side only loads and supervises; every packet-time decision happens on the DPA, inside the algorithm you write.

reactive computation. You write the algorithm in C, a special compiler (dpacc) turns it into a DPA program, and the host loader ships it onto the NIC.

The one thing your algorithm actually outputs: a rate

Your algorithm’s whole job is to set **one number per flow**: a target send **rate**. It never touches a packet and never paces anything itself. It writes the rate into a `results` struct, and when it returns, the NIC programs that flow’s **hardware rate limiter** to that value. From then on the *hardware* paces the packets — at line speed, with your algorithm nowhere in the loop — until the next event, when your code runs again to revise the number.

INFO — rate is a fraction, not a bits-per-second. The rate is a fixed-point number where $1 \ll 20$ (= 1,048,576) means “100% of line rate.” So 524288 is 50%, and a small floor value is a near-stop. You will see constants like this in the code; just read $1 \ll 20$ as “full speed.”

This split is the key idea of PCC: **you write the *policy* (what the rate should be), the NIC hardware does the *work* (pacing every packet).**

Build it and run it once

You now know the two halves and what the algorithm outputs. Before we open the algorithm itself, **build the whole thing and run it once** — to prove your toolchain works and to *see* the two halves in action: the host loader compiling your DPA program, uploading it to the NIC, and running it live.

NOTE — the algorithm is deliberately incomplete right now. The controller you’re about to run is a **scaffold**: the two reactions that move the rate are left blank (you write them in Part C). So it loads, runs, and *sees* congestion — but it won’t change the rate yet. That’s expected; this step is about the plumbing, not the policy.

Try it yourself! — build the controller and watch it run

Step 1 — build. On the card’s Arm cores, from the repo root:

```
cd doca-2 && meson setup build && ninja -C build
# => build/doca-pcc-ecn/doca_pcc_ecn_rp (the host loader; the DPA program is compiled into
↳ it)
```

That one step compiles the C you’ll edit in Part C into a **DPA image** and links it into the host loader. It builds fine even with the two reactions still blank.

Step 2 — load the controller on PF1, in the **foreground**, so all its output prints right here on the terminal:

```
sudo stdbuf -oL ./build/doca-pcc-ecn/doca_pcc_ecn_rp -d mlx5_1 -l 50
```

INFO — the flags. `-d mlx5_1` picks PF1 (the sender’s uplink — the “reaction point”); `-l 50`

is a chatty log level; `stdbuf -oL` keeps the output line-buffered so it appears as it happens. It runs in the **foreground** — you watch it live and stop it with **Ctrl-C** (that sends SIGINT, the graceful stop; never kill -9 it — see [Debugging Tips](#)).

It opens the PCC context, uploads your DPA program, and then waits for events. With no traffic yet it just sits there — that’s expected; a controller with nothing to react to has nothing to do.

Step 3 — give it something to see. In **other terminals** (leave the controller running), start the Part I marker (so packets are CE-marked and the receiver sends CNPs), then drive traffic with `benchmark.sh` — the loopback is already up from Part I:

```
# in a doca-2 terminal — the Part I marker: CE-mark every packet
sudo ./build/doca-flow/doca_flow_ecn_pcap -- --percent 100
# in another terminal (from the repo root) — runs server + client together, with a live
↪ throughput chart
./scripts/benchmark.sh
```

What you should see. Back in the controller’s terminal: now that traffic is flowing, its own `PURE_ECN` trace scrolls by — proof your algorithm is running on the DPA and seeing the CNPs:

```
PURE_ECN cnp=1    rate=1048576
PURE_ECN cnp=501  rate=1048576
PURE_ECN cnp=1001 rate=1048576
# it sees the CNPs, but the rate never moves — no reactions yet
```

Those lines are mixed in with the loader’s chatter; if it’s too busy to read, stop it and reload with `| grep --line-buffered PURE_ECN` appended to show only these. That flat rate is exactly the point of the checkpoint: **your code compiled, loaded onto the NIC’s DPA, and is running** — it just doesn’t *react* yet. Making it react is Part C. (1048576 is $1 \ll 20$ — full line rate, from the INFO box above.)

Step 4 — stop it. In the controller’s terminal press **Ctrl-C** (SIGINT — the graceful stop), then stop the marker and traffic in the other terminals.

Part B — How the controller reacts: a DCQCN loop

Start at `main()`: how the controller gets loaded

You ran this program in Part A; here is what it actually does, top to bottom. This is the **host half** — the loader and supervisor — and you won’t edit it, but following its flow shows exactly where *your* code (the DPA half) gets plugged in and takes over. It all lives in `host/pcc_ecn_rp.c`:

1. **Read the two flags** — `-d <device>` (which NIC, e.g. `mlx5_1`) and `-l <level>` (log verbosity).
2. **Open that device** — `open_pcc_device()` finds the IB device with that name that *supports PCC*, and opens it.
3. **Create a PCC context** on the device — `doca_pcc_create()`.
4. **Attach your algorithm** — `doca_pcc_set_app(pcc, pcc_ecn_rp_app)`. That `pcc_ecn_rp_app` is the **compiled DPA image**, built from `device/rp_main.c` + your `device/algo/rtt_template.c` — this single line is where the code you write gets loaded in.
5. **Configure it** — the DPA thread pool, the CNP probe format (plain RoCE CNP), logging/coredump.
6. **Start it** — `doca_pcc_start(pcc)` uploads your algorithm onto the NIC’s DPA and sets it running. **From this moment the DPA is in charge:** every congestion event runs your code.
7. **Supervise** — the host then just loops in `doca_pcc_wait()`, checking the process stays healthy and otherwise doing nothing. All the real per-event work is on the DPA now.

Steps 1–5 are setup, **step 6 is the handoff**, and step 7 is the host getting out of the way. Everything from here on — the part you actually write — runs on the DPA, once per event. That is what the rest of Part B is about.

It runs once per event, per flow

The DPA calls your algorithm **once per congestion-relevant event, for each flow** — not once per packet. The entry point is `doca_pcc_dev_user_algo()`; it hands your code that flow’s saved state, the event, and a `results` struct — you set the new rate by writing it into `results` (the function itself returns nothing). The events that matter here are:

- a **packet was sent** (a “TX” event),
- a **CNP arrived** — a Congestion Notification Packet, the receiver’s way of saying “I got a CE-marked packet, you are causing congestion” (this is the mark you set in Part I, echoed back).

INFO — the NIC batches events for you. At line rate the DPA could never keep up with one event per packet, so the hardware **coalesces** them: one event stands for many packets. That is what keeps the algorithm fast enough — it reacts per *batch*, while the hardware rate limiter does the per-packet work.

How an event reaches your handler

You don’t wire any of this up. The DPA calls **one fixed entry point** per event, and a small dispatch routes it to the right handler by **event type**. Here’s the path a **CNP** takes to the code you’ll write (TODO 1), and where the other events land:




Figure 2: The path a CNP takes to the handler you write. The dark path is the one traced here; the pale branches are where the other events land. Whichever handler runs, it writes `results->rate`.

So each of the two reactions below is just the body of one of those `..._handle_roce_*` functions — the dispatch that gets you there is already written.

The two reactions

Our controller is a textbook **DCQCN** loop — the classic “additive-increase / multiplicative-decrease” pattern. All of it lives in two short handlers in `rtt_template.c`, and each does one thing. **Both of them are left blank for you** — that’s the exercise — so below we describe exactly what each must do and the pieces you build it from. You’ll write them in Part C; here, just take in the shape.

1. A CNP arrives → cut the rate (multiplicative decrease). ← you write this, **TODO 1**. Congestion is happening, so the rate must come **down** by a fixed factor, and never fall below a floor. It goes in `..._handle_roce_cnp()`, at the marker **TODO 1**. The pieces, all already there for you: - `cur_rate` — the flow’s current rate (the fixed-point number from Part A); you edit it in place. - `ECN_CNP_DEC_FACTOR` — the cut factor, `×0.90` by default (a `#define` at the top of the file). - `doca_pcc_dev_fxp_mult(factor,`

rate) — multiplies that fixed-point factor into a rate for you. - MIN_RATE — the floor the rate must not drop below.

2. Packets keep flowing, no congestion → raise the rate (additive increase). ← you write this, **TODO 2**. When traffic is moving and nothing is cutting it, the rate should **drift back up** — gently, and only occasionally, so it doesn't overshoot and immediately re-trigger congestion. It goes in `..._handle_roce_tx()`, at the marker **TODO 2**. The pieces: - a static counter, so you act only every ~1000th call rather than on every single send event; - `AI >> 2` — a small step to add each time, about 1.25% of line rate; - `RATE_MAX` — the ceiling the rate must not exceed.

Those two reactions are the entire controller: it steers the rate on the ECN signal alone (the CNPs) — which is what *pure-ECN* means. Put together, they make the DCQCN **sawtooth**: every CNP knocks the rate down by 10%, and between CNPs it drifts back up ~1.25% at a time. When the network is congested, the down-cuts win and the rate settles low; when congestion clears (no more CNPs), the rate climbs back to full.

INFO — the one knob that changes its personality. The decrease factor is a single constant at the top of the file, `ECN_CNP_DEC_FACTOR`:

```
#define ECN_CNP_DEC_FACTOR (((1 << 16) * 900) / 1000) // ×0.90 per CNP; 800..995  
↪ = ×0.80..×0.995
```

900 means ×0.90. Make it smaller and each CNP cuts harder; larger and it barely reacts. You'll use this constant in **TODO 1**, and tune it at the end of Part C.

Part C — Write the two reactions, then watch it work

The controller is **almost complete**: the event loop, the logging, and the whole host loader are done. **The two reactions that actually move the rate are left for you** — the ones you studied in Part B:

- **TODO 1** — cut the rate when a CNP arrives (multiplicative decrease), in `..._handle_roce_cnp()`
- **TODO 2** — raise the rate when things are quiet (additive increase), in `..._handle_roce_tx()`

Both are marked in `device/algo/rtt_template.c`. As shipped they're empty, so the controller **builds and runs but never changes the rate** — a flow sends flat-out no matter how congested the link is. You'll add them one at a time and watch each half of the loop come alive.

The pieces you'll have running together while you test:

what	where	why
the ECN marker from Part I (<code>doca_flow_ecn_pcap</code>)	PF0 (m1x5_0)	marks packets CE, which makes the receiver send CNPs
the PCC controller (<code>doca_pcc_ecn_rp</code>)	PF1 (m1x5_1)	reacts to those CNPs by setting the sender's rate
RoCE traffic (<code>ib_write_bw</code>)	the SFs (m1x5_2/m1x5_3)	the flow whose rate you'll watch move

*Stage 1 — write **TODO 1** (the cut) and watch the rate collapse*

In the **Part A checkpoint** you saw the scaffold run and its rate stay **flat** under a flood of CNPs — it sees congestion but never reacts. Now you write the reaction that makes it react.

Try it yourself! — make the controller react to congestion

Step 1 — get the flow running, and note the baseline. Bring the same setup back up as in the **Part A checkpoint** — the Part I marker at `--percent 100` and `./scripts/benchmark.sh` — and note the baseline

throughput on its chart (~92 Gb/s, “full speed”). As you saw there, with TODO 1 still empty the rate sits flat and the BW doesn’t budge. Leave the traffic and marker running while you edit.

Step 2 — write TODO 1. Open `device/algo/rtt_template.c` and find TODO 1 in `..._handle_roce_cnp()`. Using the pieces from Part B — `doca_pcc_dev_fxp_mult()`, `ECN_CNP_DEC_FACTOR`, `cur_rate`, `MIN_RATE` — make the rate come **down** by the cut factor on each CNP, then clamp it up to the floor so it can’t go below `MIN_RATE`. It’s two lines. (Stuck? The [Check your answer](#) block below has it.)

Step 3 — rebuild, reload, and watch it collapse. Rebuild, then load the controller again in the foreground (the traffic and marker from Step 1 are still running):

```
ninja -C build
sudo stdbuf -oL ./build/doca-pcc-ecn/doca_pcc_ecn_rp -d mlx5_1 -l 50
```

NOTE — the DPA image is built at *configure* time. If an edit doesn’t seem to take effect, do a clean rebuild: `rm -rf build && meson setup build && ninja -C build`.

Now you should see two things:

1. **The throughput drops sharply.** With the controller cutting the rate on every CNP, the client’s BW average falls well below your baseline (it at least halves, usually much more). *That is congestion control happening* — the sender is throttling itself.
2. **The rate walking down in the controller’s terminal** — the multiplicative-decrease half of the loop, live (the `PURE_ECN cnp=... rate=...` lines, no longer flat):
PURE_ECN cnp=1 rate=943718
PURE_ECN cnp=501 rate=214380
PURE_ECN cnp=1001 rate=52428
...

Stop the controller with **Ctrl-C** (SIGINT — the graceful stop) when you’re done looking.

Stage 2 — write TODO 2 (the recovery) and complete the sawtooth

With only TODO 1, the rate can only ever fall: once CNPs cut it, it stays down even after congestion clears. TODO 2 adds the other half — the gentle climb back up — so the controller can *recover*.

Try it yourself! — let the rate come back

Step 1 — write TODO 2. Find TODO 2 in `..._handle_roce_tx()`. Using the pieces from Part B — a static counter, `AI >> 2`, `RATE_MAX` — every ~1000th call, add the small step to `cur_rate` and cap it at `RATE_MAX`. Acting only every ~1000th call is what keeps the increase *gentle* — leave the gate out and the rate rockets straight back up and re-triggers congestion.

Step 2 — rebuild: `ninja -C build`.

Step 3 — run, and this time make congestion come and go. Load the controller (foreground) and traffic as in Stage 1, then, partway through, **stop the marker** in its terminal so the CNPs dry up while the flow keeps running:

```
sudo pkill -INT -x doca_flow_ecn_pcap # congestion clears; no more CNPs
```

Watch the controller’s terminal: while the marker ran the rate was cut down (that’s TODO 1); once the CNPs stop, the rate **climbs back toward full** (that’s TODO 2). That rise-and-fall is the DCQCN **sawtooth** — both halves of your controller working together:

```
PURE_ECN cnp=1001 rate=52428 # cut down while marking
# ... marker stopped: no new PURE_ECN cuts, TX events raise the rate back to 1048576 ...
```

Stop with **Ctrl-C** (SIGINT — the graceful stop).

Check your answer

Show the two finished reactions

TODO 1 — in `..._handle_roce_cnp()`:

```
cur_rate = doca_pcc_dev_fxp_mult(ECN_CNP_DEC_FACTOR, cur_rate); // rate = rate * 0.90
if (cur_rate < MIN_RATE) cur_rate = MIN_RATE; // never below the floor
```

TODO 2 — in `..._handle_roce_tx()`:

```
static uint32_t g_tx_inc = 0;
if ((++g_tx_inc % 1000) == 0) { // every ~1000th send event
    cur_rate += (AI >> 2); // add ~1.25% of line rate
    if (cur_rate > RATE_MAX) cur_rate = RATE_MAX;
}
```

Or diff your file against the finished controller — it ships in the **solutions** repository at the same path:

```
diff device/algo/rtt_template.c
↪ <solutions-repo>/doca-2/doca-pcc-ecn/device/algo/rtt_template.c
```

Going further (optional) — tune how hard it cuts

Now that the loop works, change **how aggressively** it reacts. Every CNP currently cuts the rate by 10% (`ECN_CNP_DEC_FACTOR = ×0.90`). This one number sets the controller’s whole personality: a sharper cut empties the queue faster (fewer retransmissions, higher goodput, lower latency) up to a point; too sharp and you under-use the link.

Try it yourself! — sweep the cut factor

Open `device/algo/rtt_template.c` and change the 900 in `ECN_CNP_DEC_FACTOR` (near the top of the file). Try a gentler reaction first, rebuild, and re-run Stage 1:

```
#define ECN_CNP_DEC_FACTOR (((1 << 16) * 990) / 1000) // ×0.99 — barely cuts per CNP
```

With `×0.99` the rate barely comes down, the queue stays deep, and you’ll see **far more CNPs** in the controller’s `PURE_ECN` trace and lower goodput than at `×0.90`. Now try a **sharper** cut (`800 = ×0.80`) and compare — fewer CNPs, a shallower queue. The tuned sweet spot on our testbed was **×0.90**:

per-CNP cut	goodput	CNPs	latency
×0.99	68 Gb/s	many	high
×0.90	88.5 Gb/s	few	low
×0.80	86.7 Gb/s	fewest	very low

INFO — the surprising bit. The rate *set-point* averaged about the same across all of these, yet goodput ranged from 68 to 88 Gb/s. Goodput is governed by **queue depth and retransmissions**, not the average rate — which is why a *sharper* cut (shallower queue, fewer drops) can *raise* goodput. We have a script that sweeps this automatically and plots the whole curve — ask an organiser if you would like to see it.

Put 900 back when you’re done to restore the tuned controller.

What you built (and saw)

- The **two-halves** model: a host loader ships your algorithm onto the **DPA**, which runs it once per event and sets a per-flow **rate** that the NIC hardware then enforces.
 - The two reactions **you wrote** — `TODO 1` (CNP → cut $\times 0.90$) and `TODO 2` (TX → raise $\sim 1.25\%$) — which together make a complete **DCQCN** sawtooth.
 - **Congestion control happening live**: with only the cut written, a flow's throughput collapsing as the controller reacts to the CE marks you produced in Part I; with the recovery added too, the rate climbing back when congestion clears.
 - How **one constant** changes the whole behaviour, and why a sharper cut can *improve* goodput.
-

Debugging Tips

- **The rate never moves** — **PURE_ECN prints the same number on every line** → `TODO 1` is still empty (or you edited it but didn't rebuild). That flat output is the scaffold's as-shipped behaviour; write the multiplicative decrease in `..._handle_roce_cnp()` and rebuild. If instead the rate falls but never climbs back after congestion clears, `TODO 2` (the increase, in `..._handle_roce_tx()`) is the missing half.
 - **The controller exits a second or two after starting** → the firmware knob `USER_PROGRAMMABLE_CC` is not live. Check with `admin/local_scripts/check_pcc_ready.sh` (it should say ready). On the tutorial cards it's set for you; if you see `doca_pcc ... not supported`, that's the cause.
 - **No PURE_ECN lines ever appear** → no CNPs are reaching the controller. Make sure the Part I marker (`doca_flow_ecn_pcap --percent 100`) is running so packets are CE-marked, and that traffic is actually flowing. (If the controller's output is just too chatty to spot them, reload it with `| grep --line-buffered PURE_ECN` appended to see only that trace.)
 - **Stop it gracefully with Ctrl-C** in its terminal (SIGINT), or with `sudo pkill -INT -x doca_pcc_ecn_rp` from another terminal. A hard kill `-9` (or `docker rm -f`) leaves a "ghost" program loaded on the DPA that blocks the next run until a chip reset.
 - **Run it in the foreground** (as shown), so its output prints live on the terminal. Launched in the background over SSH it can appear to die after a few seconds — that's the launch, not the program; the foreground keeps it attached and visible.
 - **An edit to the algorithm didn't take effect** → the DPA image is built at *configure* time, so do a clean rebuild: `rm -rf build && meson setup build && ninja -C build`.
 - **Traffic must use -R** (`benchmark.sh` and `run_{server,client}.sh` already do). Without it, the flow isn't bound to your algorithm and your handlers never run — the rate won't move at all.
-

FAQs

What is USER_PROGRAMMABLE_CC and why is it needed? It's a NIC firmware setting that enables running your *own* congestion-control code on the DPA. Without it, `doca_pcc_ecn_rp` refuses to start. It only goes live after a full power-cycle of the card, so it's set up for you ahead of the tutorial.

Why does the controller attach to m1x5_1 (PF1), but traffic uses m1x5_2/m1x5_3 (the SFs)? Congestion control is a property of the *port/uplink*, so the controller attaches to the sender's physical function (PF1). The traffic rides on sub-functions of that port. One controller governs the flows on its port.

Why does the rate collapse so much at --percent 100? Because you're marking *every* packet, the receiver sends a constant stream of CNPs, so the controller thinks the network is maximally congested and keeps cutting. Mark a smaller fraction (`--percent 50`) for a gentler, more realistic signal.

Does my algorithm run for every packet? Is the DPA fast enough? No — the NIC **coalesces** events, so your code runs per *batch*, far below the packet rate, and the hardware rate limiter does the per-packet pacing. That two-tier split is exactly why it keeps up at line rate.

Why is the rate register not the same as the goodput I measure? The rate is a *set-point*; the

throughput you get also depends on queueing and retransmissions. A controller can hold a similar average rate yet deliver very different goodput — which is what the Stage 2 sweep shows.

That's the pair: in **Part I** you marked the congestion signal in the data plane, and here in **Part II** you wrote and tuned the controller that reacts to it on the DPA — the two halves of programmable transport on a SmartNIC.