
Part 2 — Do Congestion Control

Programmable Congestion Control with DOCA PCC

Programming SmartNICs: From Packet Processing to Programmable Transport

ACM SIGCOMM 2026

Fernando Ramos  Muhammad Shahbaz  Mina Tahmasbi Arashloo  Mario Baldi 

Francisco Pereira  Bilal Saleem  Tyler Liu  Chenxingyu Zhao 

 INESC-ID, Instituto Superior Técnico, University of Lisbon  University of Michigan  University of Waterloo
 NVIDIA  Purdue University  University of Washington

In **Part 1** ([doca-flow.pdf](#)) you programmed the NIC to *mark* congestion — you set the **CE** bit on packets in the data plane. In this part you write the other side of that story: the **congestion-control algorithm that reacts to those marks**, deciding how fast a flow is allowed to send — and you run it on a set of tiny processors *inside* the NIC called the **DPA**.

By the end you will have **written the two reactions at the heart of a reaction-point controller**, watched a flow’s send rate **collapse** the moment congestion appears and **recover** when it clears, and tuned how aggressively it reacts.

Prerequisites. (A) You have done **Part 1** ([doca-flow.pdf](#)); (B) one firmware knob, `USER_PROGRAMMABLE_CC=1`, must be live for any PCC program to start (we set this already for you).

Step 1 — Where the algorithm runs

A DOCA PCC program comes in two pieces that run in two different places:

- **The host program** is just a loader and supervisor. It uploads your compiled algorithm to the NIC, opens the PCC context, and then sits there keeping it alive and printing logs.
- **The algorithm** runs on the **DPA** — the *Data-Path Accelerator* inside the BlueField-3. That is where the C code you’ll edit actually executes.

INFO — what is the DPA? It’s a cluster of small processors on the NIC’s own data path — right by the wire, where the general-purpose Arm cores are too far away to set a per-flow rate fast enough. You write the algorithm in C, `dpacc` (a compiler specific to the DPA) compiles it, and the host loader ships it onto the NIC.

What the algorithm produces: a rate. For each flow it writes a target send rate into a `results` struct and returns. It never touches a packet or paces anything itself; the NIC loads that rate into the flow’s hardware rate limiter, which paces every packet according to it.

INFO — rate is a fraction, not a bits-per-second. The rate is a fixed-point number where $1 \ll 20$ ($= 1,048,576$) means “100% of line rate.” So 524288 is 50%, and a small floor value is a near-stop. You will see constants like this in the code; just read $1 \ll 20$ as “full speed.”

This split is the key idea of PCC: you write the *policy* (what the rate should be), the NIC hardware does the *work* (pacing every packet).

Step 2 — Build the base DPA application

The repository (/home/s26t/sigcomm26-tutorial-bluefield-participants) has one directory per DOCA release: doca-2 and doca-3. Pick the relevant one for your chosen Bluefield.

Which of the two is yours? Run this command to find out which version of DOCA is installed on your card:

```
$ cat /opt/mellanox/doca/applications/VERSION
2.9.1008
```

Only the first number matters: **2.x → use doca-2, 3.x → use doca-3**. The exercises are identical in the two versions, down to the same two TODOs, differing only in a few DOCA calls renamed between the generations, which are already written for you in each tree.

Move into your version’s directory now, and stay there for the rest of Part 2:

```
$ cd ~/sigcomm26-tutorial-bluefield-participants/doca-2 # or doca-3
```

This is the only place the version is ever named. Every command from here on is written relative to that directory, so you can paste them as they are whichever card you are on.

The PCC program lives in doca-pcc-ecn/, and its layout mirrors the two halves from Step 1 — a host/ side that loads and supervises, and a device/ side that runs on the DPA:

File	Runs on	What it does
host/pcc_ecn_rp.c	Arm	loads your algorithm and supervises it
device/rp_main.c	DPA	the entry point and event dispatch
device/algo/rtt_template.c	DPA	runs the custom CC algorithm

NOTE — you edit only one file. The only file you touch in this part is `rtt_template.c` (under `device/algo/`). Everything in it is written for you except two function bodies, marked `TODO 1` and `TODO 2`, which you fill in Step 3.

The DPA entry point: `device/rp_main.c`. It holds the callback the DPA runs on every event, which hands that event straight to your algorithm — `rtt_template_algo()`, in `device/algo/rtt_template.c`.

`rtt_template_algo()` runs once per event, per flow. It can be called by several event types; two of them matter here:

- A **CNP event** is triggered by the receipt of congestion notifications. This triggers the call of the function `rtt_template_handle_roce_cnp()` in `device/algo/rtt_template.c`. Currently it has no effect (it keeps the current rate of the sender unchanged).
- A **“TX” event** is triggered when a batch of packets is sent. This triggers the call of the function `rtt_template_handle_roce_tx()` in `device/algo/rtt_template.c`. Currently it also leaves the sender’s rate unchanged.

INFO — the NIC batches events for you. At line rate the DPA could never keep up with one event per packet, so the hardware coalesces them: one event stands for many packets. It reacts per batch, while the hardware rate limiter does the per-packet work.

Let’s build and run the PCC program in its current state.

NOTE — it does not react yet. If we ran this application as it is, it would **not** alter the sender’s rate based on the CNPs received: both reactions are still empty, and you write them in Step 3.




Figure 1: The path a CNP takes to the handler you write. The dark path is the one traced here; the pale branches are where the other events land. Whichever handler runs, it writes `results->rate`.

Try it yourself! Build the application and run.

We use meson and ninja to setup and build our applications. Both commands run from inside the version directory you moved into in Step 2:

```
# Setting up the build directory
# Unlike Part 1, we need to run this every time we edit the code.
$ meson setup --reconfigure build

# Compile the application
$ ninja -C build
```

NOTE — rebuild with --reconfigure. Everything under `device/` is compiled into the DPA image at configure time, not build time. Plain `ninja` has no dependency on those files, so on its own it will report no work to do and re-link the DPA image you built before your edit. After editing a TODO, always rebuild with `meson setup --reconfigure build && ninja -C build`.

Much like Part 1, keep the benchmark script running in one terminal throughout the tutorial:

```
# Keep it running!
$ ../scripts/benchmark.sh
```

Run the DOCA Flow application you developed in Part 1 to mark packets with congestion signals. Let's make the application mark 100% of the packets:

```
$ sudo ../build/doca-flow/doca_flow_ecn -- --percent 100
```

Finally let's run base PCC application

```
# Run the pcc application
# Sender is on PF1 so we need to run the PCC application on the sender's uplink (mlx5_1)
$ sudo stdbuf -oL ./build/doca-pcc-ecn/doca_pcc_ecn_rp -d mlx5_1 -l 50
```

What you should see.

By looking at the output of the DOCA Flow application you should see that every packet is having its CE bit set, and from the PCC output you should see that the congestion notifications are being received:

```
PURE_ECN cnp=1    rate=1048576
PURE_ECN cnp=501  rate=1048576
PURE_ECN cnp=1001 rate=1048576
```

However by looking at the benchmark.sh output you *don't* see the performance being affected despite the congestion signals.

The goal of this tutorial is to implement the custom congestion-control algorithm that reacts to the congestion signals and adjusts the sender's rate accordingly.

Step 3 — Implementing the custom CC algorithm on the DPA

The custom CC algorithm has 2 main components, TODO 1 and TODO 2, both in rtt_template.c. You need to implement:

- **A multiplicative decrease** reaction to congestion signals, so the sender decreases its rate when congestion is detected.
- **An additive increase** reaction to the absence of congestion signals, so the sender increases its rate when no congestion is detected.

We will now go over the details of each of these components and how to implement them.

Multiplicative decrease (TODO 1). As CNP events arrive, decrease the rate by a fixed factor, never falling below a floor. Implement this logic in the function rtt_template_handle_roce_cnp():

- Modify the flow's current rate by setting the cur_rate variable (input parameter).
- Use ECN_CNP_DEC_FACTOR as the multiplicative decrease factor.
- Use doca_pcc_dev_fxp_mult() to perform a fast fixed-point multiplication of the rate by that factor.
- Use MIN_RATE to prevent the rate from dropping below a certain value.

INFO — the one knob that changes its behavior. The decrease factor is a single constant at the top of rtt_template.c:

```
/*×0.90 per CNP; 800..995 = ×0.80..×0.995
/*`900` means ×0.90. Make it smaller and each CNP cuts harder; larger and it
↪ barely reacts.
#define ECN_CNP_DEC_FACTOR (((1 << 16) * 900) / 1000)
```

Try it now

Rerun the experiment with DOCA Flow and benchmark.sh running on other terminals (*remember to reconfigure and rebuild the PCC application with meson setup --reconfigure build && ninja -C*

build). See now on the `benchmark.sh` output that the congestion signals *deeply* affect throughput, which suddenly plummeted when the congestion signals arrived.

Additive increase (TODO 2). When traffic is flowing and no congestion is detected, the rate should **drift back up** gently so that we don't cause congestion again. Implement this logic in the function `rtt_template_handle_roce_tx()`:

- Keep a global counter (that persists across calls) so you act only every ~1000th call rather than on every single send event;
- Increment the sender's rate by `AI >> 2`. `AI` is a predefined constant equal to 5% of line rate — it lives in `rtt_template_algo_params.h`, beside the file you are editing — so `AI >> 2`, a quarter of it, adds a gentle ~1.25% each time;
- Prevent the rate from exceeding `RATE_MAX`.

Try to run it again, now with the additive increase implemented (remember to reconfigure and rebuild the PCC application with `meson setup --reconfigure build && ninja -C build`).

Rerun the experiment with DOCA Flow and `benchmark.sh` running on other terminals. See now on the `benchmark.sh` output that the congestion signals still *deeply* affect throughput.

Try running it now with other `--percent` values for the DOCA Flow application (try it with `--percent 0.1` and `--percent 0.01`). See how the sender converged to a higher rate, depending on how much congestion was detected!

Ask an organiser if you're stuck. That's what we're here for.

Debugging tips

- **An edit didn't take effect and ninja said no work to do** → you rebuilt with plain `ninja`. The DPA image is built at *configure* time, and nothing under `device/` is in `ninja`'s dependency graph, so it cannot see your edit. Rebuild with `meson setup --reconfigure build && ninja -C build`. (`rm -rf build && meson setup build` also works and is what to reach for if the build itself looks confused, but it is far slower.)
- **No PURE_ECN lines ever appear** → no CNPs are reaching the controller. Make sure the Part 1 marker (`doca_flow_ecn --percent 100`) is running so packets are CE-marked, and that traffic is actually flowing.
- **Traffic must use -R** (`benchmark.sh` already does). Without it the flow isn't bound to your algorithm and your handlers never run — the rate won't move at all.
- **Run it in the foreground** (as shown), so its output prints live. Launched in the background over SSH it can appear to die after a few seconds — that's the launch, not the program.

Appendix A — The DPA and the two halves, in more detail

Reference material. You do not need any of this to finish the exercise.

Why two halves, and why the DPA. The Arm cores are general-purpose Linux CPUs, too far from the wire to make a per-flow rate decision fast enough. The **DPA** (Data-Path Accelerator) is a cluster of many-threaded RISC-V processors sitting on the data path, built for tiny, frequent, reactive computation. So the algorithm runs there, and the host program on the Arm is only a loader and supervisor: it compiles the DPA image (dpacc), uploads it, opens the PCC context, and then keeps the process alive. Nothing on the host runs per event.

USER_PROGRAMMABLE_CC. Running your *own* congestion-control code on the DPA is gated by a NIC firmware setting, `USER_PROGRAMMABLE_CC=1`. Without it, `doca_pcc_ecn_rp` refuses to start (it exits a second or two in). It only takes effect after a full power-cycle of the card, so it's set up for you ahead of the tutorial. Check it with `admin/local_scripts/check_pcc_ready.sh` — it should say ready.

Why the controller attaches to PF1 but traffic uses the SFs. Congestion control is a property of the *port/uplink*, so the controller attaches to the sender's physical function, PF1 (`m1x5_1`). The traffic rides on sub-functions of that port (`m1x5_2/m1x5_3`). One controller governs all the flows on its port.

Appendix B — The controller's model in full

Reference material. Step 3 has the working subset.

Rate is a fixed-point fraction. The rate is not bits per second; it is a fixed-point number where $1 \ll 20$ (1,048,576) is 100% of line rate. 524288 is 50%, a small floor value is a near-stop. Your algorithm writes this number into `results->rate`; the NIC's hardware rate limiter enforces it until your code runs again.

Events are coalesced. Your algorithm does **not** run per packet — at line rate the DPA could never keep up. The hardware batches events so one event stands for many packets, and your handler runs per *batch* while the rate limiter does the per-packet pacing. That two-tier split is exactly why it keeps up at line rate.

Why the rate collapses so hard at --percent 100. Marking *every* packet makes the receiver send a constant stream of CNPs, so the controller thinks the link is maximally congested and keeps cutting. A much smaller fraction — the `--percent 0.1` and `--percent 0.01` you try in Step 3 — is a gentler, more realistic signal.

Rate set-point vs measured goodput. The rate is a *set-point*; the throughput you actually get also depends on queueing and retransmissions. A controller can hold a similar average rate yet deliver very different goodput: a sharper cut keeps the queue shallow and can *raise* goodput even though the average rate is unchanged.

Appendix C — The code you edit, and the API

The DPA headers on the card are the authority. This is only a map of what matters here.

The file you edit is `device/algo/rtt_template.c`. Its two handlers are the whole exercise:

Function	Event	You write
<code>..._handle_roce_cnp()</code>	a CNP	TODO 1 — multiplicative decrease (cut $\times 0.90$)
<code>..._handle_roce_tx()</code>	a send	TODO 2 — additive increase (raise $\sim 1.25\%$)

The dispatch that routes an event to the right handler, and the entry point `doca_pcc_dev_user_algo()` (in `device/rp_main.c`), are already written for you.

The DPA helpers and constants you use:

Name	What it is
<code>cur_rate</code>	the flow's current rate (fixed-point); edit it in place
<code>doca_pcc_dev_fxp_mult(f, r)</code>	multiply a fixed-point factor <code>f</code> into a rate <code>r</code>
<code>ECN_CNP_DEC_FACTOR</code>	the per-CNP cut factor ($\times 0.90$ by default); the one tuning knob
<code>MIN_RATE / RATE_MAX</code>	the floor and ceiling the rate must stay within
<code>AI</code>	the additive-increase step; <code>AI >> 2</code> is $\sim 1.25\%$ of line rate

The host loader (`host/pcc_ecn_rp.c`, which you don't edit) uses `doca_pcc_create`, `doca_pcc_set_app`, `doca_pcc_start`, and `doca_pcc_wait` — open the context, attach the DPA image, start it, supervise.

The program's own flags: `-d <device>` (which NIC, e.g. `mlx5_1`) and `-l <level>` (log verbosity).

Further reading

- [DOCA PCC programming guide](#) — search “DOCA PCC”; swap the archive version to match your card.
- [DOCA DPA subsystem](#) — how DPA programs are built and loaded.