
Bonus — Path Steering

Combining DOCA PCC with DOCA Flow

Programming SmartNICs: From Packet Processing to Programmable Transport

ACM SIGCOMM 2026

Fernando Ramos  Muhammad Shahbaz  Mina Tahmasbi Arashloo  Mario Baldi 

Francisco Pereira  Bilal Saleem  Tyler Liu  Chenxingyu Zhao 

 INESC-ID, Instituto Superior Técnico, University of Lisbon  University of Michigan  University of Waterloo
 NVIDIA  Purdue University  University of Washington

Earlier parts of the tutorial introduced two complementary BlueField programming models:

- **DOCA Flow** programs packet classification, modification, forwarding, and hardware counters in the NIC data plane.
- **DOCA PCC** runs a congestion-control algorithm on the DPA and reacts to RoCE transport events for individual queue pairs (QPs).

This bonus section combines them into a closed-loop, two-path steering system. PCC continues to calculate a congestion-derived value for each QP, but it does not apply that value as a NIC rate limit. Instead, the DPA reports it to the host, where it becomes an input to the DOCA Flow controller. The controller converts the relative congestion signals into a packet-splitting ratio and updates an existing hardware dispatch table while traffic is running.

The experiment runs on one BlueField-3 using a physical port loopback. Two virtual paths share the same cable and are distinguished by a temporary DSCP bit. Two parallel RDMA QPs act as path-specific congestion probes. The sender pipeline randomly assigns packets to paths; the receiver pipeline emulates a different ECN policy for each path, removes the private marker, and delivers the two QPs to separate receiver SFs.

By the end of this section, you will understand how to:

1. preserve path identity when standard RoCE CNP feedback identifies only a QP;
2. build a random HASH classifier whose distribution can be changed without recreating the HASH pipe;
3. connect per-QP PCC reports to live DOCA Flow entry updates; and
4. validate the complete data and feedback loop using hardware counters and host-side diagnostics.

The implementation supports DOCA 2.9, 3.1, and 3.4 through a shared compatibility layer. The 3.1 and 3.4 paths are hardware-validated; the 2.9 path is compile-tested and still needs hardware validation. DOCA 2.7 is not supported. The logical pipeline and command-line model remain the same; only the SDK-specific host, PCC-reporting, and DOCA Flow APIs differ.

Step 1 — Understand the scenario

Step 1.1 — Why a single QP is not enough

Start with the intuitive design: one sender, one receiver, and a path multiplexer that distributes packets from one RDMA queue pair (QP) over two paths.




Figure 1: A single-flow multi-path scenario

The multiplexer can choose a path independently for every packet, but ordinary RoCE congestion feedback does not preserve that choice. A receiver that sees an ECN-marked packet sends a Congestion Notification Packet (CNP) for the affected QP. When the CNP reaches the sender, PCC can identify the QP, but it cannot tell whether the original data packet traversed path 0 or path 1. Both paths have collapsed into one transport-level feedback stream.

This creates an observability problem. Suppose path 0 is congested and path 1 is clear. The sender may learn that the QP should slow down, but it has no path-specific signal from which to calculate a better packet-splitting ratio. Changing the QP rate alone also reduces traffic on both paths, including the uncongested one.

The tutorial therefore separates two decisions:

1. PCC estimates congestion from transport feedback and exports a rate-like signal for each QP.
2. DOCA Flow uses those signals to change the fraction of packets assigned to each path.

PCC does not enforce the calculated rate in this example. The DPA program returns `DOCA_PCC_DEV_MAX_RATE` to the NIC and reports its calculated value to the host as a steering signal. DOCA Flow, rather than the normal PCC rate limiter, controls the traffic distribution.

Step 1.2 — Use paired QPs as path probes

To recover path identity, the application uses two parallel RDMA QPs with equal or comparable offered load. We call them the blue QP and the green QP.

The colors identify QPs, not paths. Packets from either QP may be assigned to either virtual path by the sender-side multiplexer. At the receiver-side path emulator, however, each path marks only its designated probe QP:

Selected virtual path	QP eligible for ECN marking	Interpretation
Path 0	Blue QP	Blue-QP congestion represents path 0
Path 1	Green QP	Green-QP congestion represents path 1

Packets from the other QP still traverse and consume capacity on that path; they are simply not used as




Figure 2: A two-flow multi-path scenario

that path’s congestion probe. Because only roughly half of a paired workload is eligible for marking on a given path, the receiver-side sampler uses twice the requested all-traffic marking probability, capped at 100%. For example, an intended 5% path marking rate becomes a 10% sampling probability for the designated QP.

This compensation assumes the paired QPs contribute similar traffic volumes. If one QP is idle or substantially slower, its PCC signal is no longer a representative probe for the corresponding path. That is a property of this tutorial construction, not a general requirement of multipath steering.

The two QPs use different receiver destination IP addresses. The destination IP identifies blue versus green traffic and lets the receiver deliver each flow to its receiver SF. It does **not** encode the packet's selected path; that is a separate per-packet marker.


Step 1.3 – Map the scenario onto one BlueField-3

The tutorial emulates two paths using the two PF domains and one physical loopback cable on a single BlueField-3:

- The sender application uses a sender SF in the PF1 domain.
- PF1 runs the egress steering pipeline and sends traffic through physical port p1.
- A loopback cable carries the traffic from p1 to p0.
- PF0 runs the ingress path-emulation pipeline.
- Two receiver SFs represent the blue and green receiver endpoints.

The earlier tutorial setup already introduced `ns0 (mlx5_2, 10.0.0.1)` and `ns1 (mlx5_3)`. This bonus setup adds one more receiver endpoint: `pf0sf4` appears as RDMA device `mlx5_4` inside namespace `ns0_1`, with IP address `10.0.0.11`. It is a second SF in the PF0 domain, not a second physical port. The tutorial card is prepared with this extra SF and namespace before you begin.

Endpoint	Representor	RDMA device	Namespace	IP	Role
Blue receiver SF	pf0sf0	mlx5_2	ns0	10.0.0.1	Blue QP server
Green receiver SF	pf0sf4	mlx5_4	ns0_1	10.0.0.11	Green QP server
Shared sender SF	pf1sf0	mlx5_3	ns1	10.0.0.2	Both QP clients



and clears only that private bit before delivering the packet. Existing ECN bits and unrelated DSCP bits are preserved.

The UDP destination port is deliberately not used as a path marker. UDP port 4791 is covered by the RoCEv2 invariant CRC; rewriting it causes the receiver to drop the packet even if a later rule restores the original value.

Step 1.4 — Follow a packet through the forward path

For each outgoing RoCEv2 packet, the data plane performs the following steps:

1. The sender posts traffic on the blue or green RDMA QP through the sender SF.
2. PF1 admits RoCEv2 traffic to EGRESS_CLASSIFY. A 64-entry random HASH pipe selects a persistent bucket for the packet.
3. A metadata dispatch table maps that bucket to path 0 or path 1. Changing the number of buckets assigned to each path changes the steering ratio without rebuilding the HASH pipe.
4. The selected path-rewrite pipe writes DSCP path bit 0 or 1, and the packet is sent through the physical loopback.
5. PF0 reads the path bit first, then checks the destination IP. Only the probe QP designated for that path enters the path's probabilistic CE-marking stage.
6. PF0 clears the private path bit and uses the destination IP to deliver blue traffic to receiver SF0 and green traffic to receiver SF1.

Non-RoCE traffic bypasses the random classifier.

Step 1.5 — Close the feedback and steering loop

Path selection is per packet, while congestion control and QP identification remain per transport flow. The feedback and steering loop operates as follows:

1. PCC calculates a congestion-derived value for each sender QP on the DPA using the same algorithm introduced in the DOCA PCC portion of the tutorial.
2. PCC reports these values to the host through binary trace reports on DOCA 3.x or a periodically polled PCC mailbox on DOCA 2.9.
3. The host averages reports over the interval and applies a persistent EWMA. Flows that have not yet been associated with a path are reported as pending-map and do not influence the ratio.
4. The controller compares the aggregate reduced-rate signals for the two path groups and converts the result into a number of path-0 buckets out of 64.
5. Only dispatch entries whose assignment crosses the old/new boundary are updated. No HASH entries are added or removed while traffic is running.

At startup, 32 buckets are assigned to each path. The controller retains at least three buckets for a constrained path in the one-sided full-rate case, so both paths continue receiving probe traffic. The 64-bucket representation gives a steering granularity of $1/64$, or approximately 1.56%.

Step 1.6 — What the experiment demonstrates

This experiment demonstrates how PCC and DOCA Flow can form a closed-loop traffic-steering system. PCC measures congestion for the probe QPs, and the host converts those measurements into a desired path share. DOCA Flow applies that share by updating the bucket dispatch entries while traffic continues, without rebuilding the random HASH pipe. Together, PCC provides the feedback signal and DOCA Flow enforces the resulting per-packet steering decision at line rate. The single-card loopback emulates the two paths for the tutorial, but the same control and steering design can target two real next hops.

Step 2 — Run and observe the completed solution

Before writing any DOCA Flow code, run the completed implementation and observe the behavior that your implementation must reproduce. This gives you a known-good reference for the pipeline startup messages, initial path split, and hardware counters.

Step 2.1 — Build the solution

Enter the path-steering project, configure its build directory, and compile the PCC and steering applications:

```
$ cd doca-3/pcc-path-steering
$ meson setup build
$ ninja -C build
```

If build is already configured, use `meson setup build --reconfigure` instead of the first-time `meson setup build`. Run all remaining relative commands in this guide from the `doca-3/pcc-path-steering` directory.

The same source builds on DOCA 2.9, 3.1, and 3.4 through the compatibility layer. This checkout keeps the shared source under `doca-3/pcc-path-steering`; there is no separate `doca-2/pcc-path-steering` directory.

Step 2.2 — Start the receiver-side path emulator

In the first terminal, start the ingress role on PF0 with the two receiver SF representors:

```
sudo ./build/doca_flow_steer \
-r pci/0000:03:00.0,pf0sf0 \
-R pci/0000:03:00.0,pf0sf4 \
--path0-ip 10.0.0.1 \
--path1-ip 10.0.0.11 \
--role ingress \
--path0-percent 0.025 \
--path1-percent 0.05
```

This process reads the private path bit, applies the configured path-specific ECN policy, clears the private bit, and delivers each QP to its receiver SF.

The values are literal percentages, not fractions: `0.025` means **0.025%**, and `0.05` means **0.05%**. The selected-probe sampler doubles them internally to 0.05% and 0.1%, respectively, then rounds down to a supported power-of-two probability. These low rates still produce feedback quickly at line rate while making path 1 approximately twice as likely to mark its probe QP.

Step 2.3 — Start PCC with the completed egress solution

In a second terminal, start PCC on the sender and pass the PF1 sender SF representor to enable its embedded egress steering pipeline:

```
sudo ./build/doca_pcc \
--device mlx5_1 \
-r pci/0000:03:00.1,pf1sf0 \
--path0-ip 10.0.0.1 \
--path1-ip 10.0.0.11
```

On DOCA 3.x, successful setup includes these messages:

```
EGRESS_BUCKET_DISPATCH ready: meta.u32[4] bucket -> changeable path forward
Classify random-bucket pipe ready: HASH/random writes bucket to meta.u32[4]
```

On DOCA 2.9, the classifier message is instead:

```
Legacy random HASH ready: 6 bits, 64 immutable metadata buckets
```

The dispatch-ready message appears before the classifier-ready message on all supported versions.

Step 2.4 — Generate traffic and observe the result

Start an `ib_write_bw` receiver for the first flow in namespace `ns0`:

```
sudo ip netns exec ns0 ib_write_bw \
  -d mlx5_2 -R --report_gbits --run_infinitely &
```

Start the receiver for the second flow in namespace `ns0_1`:

```
sudo ip netns exec ns0_1 ib_write_bw \
  -d mlx5_4 -R --report_gbits --run_infinitely &
```

Run the first sender flow from namespace `ns1` in a new terminal:

```
sudo ip netns exec ns1 ib_write_bw \
  -d mlx5_3 -R -F 10.0.0.1 --report_gbits --run_infinitely
```

Run the second sender flow from the same namespace in another terminal so that both flows remain active concurrently:

```
sudo ip netns exec ns1 ib_write_bw \
  -d mlx5_3 -R -F 10.0.0.11 --report_gbits --run_infinitely
```

The PCC process should report packets assigned to both paths with the initial 32:32 bucket split:

```
egress assigned counters: path0=<count> path1=<count> ratio=32:32 buckets
```

The two packet counts will not be exactly equal because bucket selection is random. Before PCC reacts to congestion, however, they should show that the initial bucket assignment sends traffic over both paths at roughly the same rate.

Within the first few reporting intervals, the PCC terminal should also show one mapping for each sender QP, for example:

```
RDMA-CM REQ grouping: sender QPN 0x<qpn> receiver IP=0x<ip> -> path0
RDMA-CM REQ grouping: sender QPN 0x<qpn> receiver IP=0x<ip> -> path1
```

The periodic PCC path-share diagnostic should then report `pending-map=0`. Mapping and a share change should normally complete within a few seconds. If `pending-map` remains nonzero, verify that both clients use `-R`, both receiver processes were started first, and both destination IPs match the ingress and PCC commands.

You should be able to observe the following:

- Both `ib_write_bw` clients continue reporting throughput without transport errors, and the two parallel flows achieve similar throughput.

- Path 1 produces a stronger congestion signal because its configured ECN-marking probability is twice that of path 0.
- The raw CE rates may not have an exact 2:1 ratio because PCC changes how much traffic each path carries.
- Dividing each path's CE rate by its throughput should show that path 1 marks packets at approximately twice the rate of path 0.
- PCC eventually applies more than 32 of the 64 buckets to path 0 because path 1 has the higher configured ECN probability.
- After the share changes, path 0 carries more traffic than path 1 because more buckets now forward to path 0.

The ingress process prints the path throughput and new CE marks once per polling interval:

```
ingress throughput: path0=<Gbps> Gbps path1=<Gbps> Gbps total=<Gbps> Gbps
ingress newly CE-marked: path0=<total> (<delta>, <rate>/s) path1=<total> (<delta>, <rate>/s)
```

Compare CE deltas only over the same measurement window as the path-throughput samples. Raw CE counts alone can be misleading after PCC changes the amount of traffic assigned to each path.

When PCC reacts to the higher normalized congestion signal on path 1, it prints a share update whose new path-0 value is greater than 32:

```
PCC path share applied: path0 32-><new-path0-share>/64 path1=<new-path1-share>/64
```

The later egress counter messages should show the newly applied bucket ratio:

```
egress assigned counters: path0=<count> path1=<count> ratio=<path0>:<path1> buckets
```

Keep this output available as a reference when you implement the classifier and dispatch logic in the following steps. Stop the four `ib_write_bw` processes, PCC, and the ingress steering process before moving to the DIY implementation.

Step 3 — Randomly classify packets

Before changing the path split, we first need to assign packets to a stable set of buckets that the controller can assign to either path. This step walks through the provided random classifier in [steering/steer.c](#) and then asks you to build the same kind of pipe yourself.

Step 3.1 — Understand the random-bucket classifier

The classifier maps every admitted RoCEv2 packet to one of 64 buckets. The selected bucket number is written to `meta.u32[4]`, and the packet is forwarded to `EGRESS_BUCKET_DISPATCH`.

```
RoCEv2 packet
|
v
EGRESS_CLASSIFY (RANDOM HASH, 64 entries)
|
| write bucket 0..63 to meta.u32[4]
v
EGRESS_BUCKET_DISPATCH
```

The HASH entries never need to change after they are installed. This is important because the random assignment mechanism and the path policy have different lifetimes. The HASH pipe supplies a persistent set of buckets, while the following dispatch pipe decides which path currently owns each bucket.

The per-packet application metadata word `meta.u32[4]` is used intentionally to carry the selected bucket index to the dispatch pipe. DOCA Flow HASH pipes use part of `u32[3]` internally, so the application stores its bucket number outside that region.

Step 3.2 — Understand the DOCA 3.x implementation

DOCA 3.x provides a native HASH-pipe RANDOM map algorithm. `create_classify_pipe()` selects it with `DOCA_FLOW_PIPE_HASH_MAP_ALGORITHM_RANDOM`, so the pipe assigns each packet to one of its HASH entries without requiring a packet field as the hash key. The action template allows each HASH entry to write a value to `meta.u32[4]`. The pipe's fixed forward sends every selected entry to the dispatch pipe.

The important pipe properties are:

- The pipe type is `DOCA_FLOW_PIPE_HASH`.
- The map algorithm is `DOCA_FLOW_PIPE_HASH_MAP_ALGORITHM_RANDOM`.
- The pipe contains `PATH_SHARE_BUCKETS` entries, which is 64 entries in this tutorial.
- The action template makes `meta.u32[4]` writable.
- The pipe forwards to `EGRESS_BUCKET_DISPATCH`.

After creating the pipe, `add_classify_entries()` installs the 64 immutable entries as one batch. Entry `idx` writes `RTE_BE32(idx)` to `meta.u32[4]`. The first 63 entries use `STEER_WAIT_FOR_BATCH`, and the final entry submits the batch for processing.

Step 3.3 — Understand the DOCA 2.9 implementation

DOCA 2.9 does not provide the native RANDOM map algorithm used on DOCA 3.x. Instead, the packet parser supplies a per-packet random value in `parser_meta.random`, and `create_legacy_small_random_table()` configures a HASH pipe to hash that value. The match mask sets all 16 bits of `parser_meta.random`, allowing the HASH pipe to distribute packets across its 64 entries. Each entry writes its bucket index to `meta.u32[4]` and forwards to the same dispatch pipe used by the DOCA 3.x implementation. The rest of the steering pipeline therefore sees the same metadata contract on every supported DOCA version.

The compatibility choice is selected at compile time through `STEER_USE_RANDOM_HASH_CLASSIFIER` and `DOCA_USES_LEGACY_FLOW_BACKEND` in [steering/doca_flow_compat.h](#). The tutorial code should use the compatibility wrappers such as `steer_pipe_hash_add_entry()` instead of duplicating SDK-specific entry API calls.

Step 3.4 — Build the classifier

In Step 2, the completed solution created the random HASH pipe and reported traffic on both paths with an initial 32:32 bucket split. Your implementation should reproduce the same startup messages and counter behavior. For now, however, it will maintain a fixed 32:32 bucket-to-path split; you will add dynamic steering in the next section.

Try it yourself! Build the random-bucket classifier.

The repository contains the completed implementation so that you can run the working example first. To replace the relevant solution bodies with numbered TODO stubs, apply the provided patch from its project directory:

```
$ cd doca-3/pcc-path-steering
$ patch -p1 < solution-to-diy.patch
```

The patch exposes all four TODOs at once. In this step, leave **TODO 4 unfinished** so `apply_path_share()` remains a no-op; you will implement it in Step 4.

Implement a 64-entry random-bucket classifier that provides the same metadata contract as the walkthrough above. The TODO comments point back to the relevant guide sections:

- **DOCA 2.9:** implement TODO 1 in `create_legacy_small_random_table()` using Step 3.3.
- **DOCA 3.1 or 3.4:** implement TODO 2 in `create_classify_pipe()` and TODO 3 in `add_classify_entries()` using Step 3.2.

Only the branch for the installed SDK is compiled, so you do not need to implement the other version's classifier during the tutorial. The DIY version keeps the completed dispatch pipe and assigns 32 buckets to each path. Its `apply_path_share()` function is initially a no-op, so PCC diagnostics may calculate a new target while the applied hardware ratio remains 32:32.

On DOCA 3.1 or 3.4, your implementation should:

1. create a non-root HASH pipe named `EGRESS_CLASSIFY`;
2. configure the native RANDOM map algorithm when `STEER_USE_RANDOM_HASH_CLASSIFIER` is enabled;
3. make `meta.u32[4]` writable through the action template;
4. forward the HASH pipe to the provided dispatch pipe;
5. add one entry for every bucket from 0 through 63;
6. write the entry's bucket number to `meta.u32[4]` in big-endian order; and
7. submit the entries as a batch and verify all 64 completions.

For the DOCA 2.9 branch, create the equivalent immutable HASH table with `create_legacy_small_random_table()` while preserving the same metadata output. Use `steer_pipe_hash_add_entry()` so the compatibility layer handles the SDK-specific function signature.

Rebuild from the project directory:

```
$ ninja -C build
```

Then repeat Steps 2.2 through 2.4: start ingress, start PCC, start both receivers, and start both clients. You are finished when the application prints the classifier-ready message for your SDK and reports traffic on both paths with a fixed 32:32 bucket split. The path-share diagnostic may calculate a target other than 32:32, but the applied ratio must remain 32:32 because TODO 4 is intentionally unfinished. Stop all six processes before continuing to Step 4.

Step 4 — Steer packets dynamically

The random classifier decides only which bucket a packet belongs to. The DIY starting point already provides the metadata dispatch table with a fixed 32:32 assignment so that you can test Step 3 independently. This step explains that table and asks you to implement the live path-share update.

Step 4.1 — Separate classification from path selection

`create_classify_dispatch_pipe()` creates a BASIC pipe named `EGRESS_BUCKET_DISPATCH`. The pipe matches the bucket number in `meta.u32[4]` and uses a changeable forward for each entry.

```
meta.u32[4] bucket
    |
    v
EGRESS_BUCKET_DISPATCH
    |
```

```
+---- buckets assigned to path 0 ----> EGRESS_PATH0_REWRITE
|
+---- buckets assigned to path 1 ----> EGRESS_PATH1_REWRITE
```

The dispatch pipe installs 64 live entries but reserves capacity for 128 entries. The extra capacity is required because the hardware steering backend can temporarily allocate a replacement rule while an entry is being updated.

At startup, buckets 0 through 31 forward to EGRESS_PATH0_REWRITE, and buckets 32 through 63 forward to EGRESS_PATH1_REWRITE. The two rewrite pipes set the private DSCP path bit and then forward the packet toward the wire.

Step 4.2 — Update the path share

`apply_path_share()` receives the desired number of path-0 buckets as an integer from 0 through 64. A bucket belongs to path 0 when its index is less than `path0_share`, and it belongs to path 1 otherwise.

For example, changing the split from 32:32 to 48:16 moves only buckets 32 through 47 from path 1 to path 0. The other 48 dispatch entries already have the desired forward and are left untouched.

For every bucket that crosses the boundary, the function:

1. creates a DOCA_FLOW_FWD_PIPE forward to the desired path-rewrite pipe;
2. calls `steer_pipe_update_entry()` on the existing dispatch entry;
3. processes and checks the update completion;
4. records the new bucket-to-path assignment only after the update succeeds; and
5. drains pending replacement operations after all changed buckets are processed.

The HASH entries are not modified, removed, or re-created during this operation. Only the changeable forwards in the BASIC dispatch pipe are updated.

Step 4.3 — See dynamic steering in action

Run the ingress path emulator as described in Step 2.2, start the PCC application with embedded egress steering, and then generate traffic on both QPs. When the two path signals request a different split, the controller prints the old and new assignment:

```
PCC path share applied: path0 32->48/64 path1=16/64
```

The periodic hardware counters should subsequently report the new bucket ratio:

```
egress assigned counters: path0=<count> path1=<count> ratio=48:16 buckets
```

The packet-counter ratio will approach the configured bucket ratio over time rather than changing to an exact value immediately. The update log proves that the dispatch policy changed, while the counters prove that later packets followed the new policy.

Try it yourself! Implement live path-share updates.

If you have not already done so in Step 3, enter `doca-3/pcc-path-steering` and apply `solution-to-diy.patch` there to expose the numbered TODO stubs. TODO 4 in `apply_path_share()` points back to Step 4.2.

The provided `create_classify_dispatch_pipe()` already installs the 64-entry dispatch table, saves the entry handles, and records the initial path assignment. Implement TODO 4 in `apply_path_share()` so that it:

1. validates the requested path-0 share;
2. returns immediately when the requested share is already active;
3. updates only buckets whose desired path differs from their recorded path;
4. waits for and validates every entry-update completion;
5. leaves the recorded assignment unchanged when an update fails;
6. drains the replacement operations before returning; and
7. updates `applied_path0_share` only after all required bucket updates succeed.

Use the compatibility wrapper `steer_pipe_update_entry()` rather than calling a version-specific DOCA Flow update API directly.

As a quick reasoning check, verify that a transition from 32:32 to 48:16 performs exactly 16 entry updates and that repeating 48:16 performs none. Rebuild with `ninja -C build`, then repeat Steps 2.2 through 2.4. You are finished when traffic continues during a share change, the application prints the applied-share message within a few reporting intervals, and the hardware counters move toward the new ratio.

FAQ

Why not steer traffic dynamically with the HASH pipe itself?

Many DOCA versions do not support updating HASH-pipe entries after they have been created. The tutorial therefore keeps the HASH pipe immutable and performs dynamic steering in a separate BASIC dispatch pipe whose entries can be updated while traffic is running.

Can I use a different number of buckets?

Yes, but the supported HASH-pipe size depends on the DOCA version. You may change to a bucket count that stays within the limit of every DOCA version you intend to support, and update the initial split and related constants accordingly.